

1. Create a schema based on the given dataset

```
Time taken: 0.012 seconds
hive> create table LoginReport(
  >     SL_NO int,
  >     Agent string,
  >     Date_l string,
  >     Login_time string,
  >     logout_time string,
  >     Duration string
  > )
  > row format delimited
  > fields terminated by ',';
OK
Time taken: 0.077 seconds
hive> create table Performance(
  >     SL_NO int,
  >     Date_l string,
  >     Agent_Name string,
  >     Total_Chats int,
  >     Average_Response_Time string,
  >     Average_Resolution_Time string,
  >     Average_Rating float,
  >     Total_Feedback int
  > )
  > row format delimited
  > fields terminated by ',';
OK
Time taken: 0.08 seconds
hive> █
```

2. Dump the data inside the hdfs in the given schema location.

```
Time taken: 0.069 seconds
hive> load data local inpath 'file:///tmp/AgentLoggingRe
FAILED: SemanticException Line 1:23 Invalid path ''file
hive> load data local inpath 'file:///tmp/data/AgentLo
Loading data to table db1.loginreport
Table db1.loginreport stats: [numFiles=1, totalSize=55
OK
Time taken: 0.669 seconds
hive> load data local inpath 'file:///tmp/data/AgentPe
Loading data to table db1.performance
Table db1.performance stats: [numFiles=1, totalSize=10
OK
Time taken: 0.241 seconds
hive> █
```

3. List of all agents' names.

```
hive> select Agent from LoginReport;  
OK  
Agent  
Shivananda Sonwane  
Khushboo Priya  
Nandani Gupta  
Hrisikesh Neogi  
Mukesh  
Sowmiya Sivakumar  
Manjunatha A  
Harikrishnan Shaji  
Suraj S Bilgi  
Shivan K  
Anurag Tiwari  
Ishawant Kumar  
Shivan K  
Shubham Sharma  
Shivan K  
Prerna Singh  
Shivan K  
Shivan K  
Hrisikesh Neogi  
Khushboo Priya  
Nishtha Jain  
Prateek _iot  
Hrisikesh Neogi  
Shivan K  
Nishtha Jain  
Khushboo Priya  
Suraj S Bilgi  
Mithun S  
Harikrishnan Shaji  
Manjunatha A  
Hrisikesh Neogi  
Sowmiya Sivakumar  
Madhulika G  
Boktiar Ahmed Bappy
```

4. Find out agent average rating.

```
hive> select Average_Rating from Performance limit 20;
OK
NULL
4.11
3.14
4.55
4.71
3.67
4.62
5.0
5.0
4.38
4.12
3.6
4.43
4.36
4.57
4.3
4.17
4.75
4.25
4.05
Time taken: 0.06 seconds, Fetched: 20 row(s)
hive> █
```

5. **Total working days for each agents**

```
FAILED: ParseException line 1:03 missing EOF at 'by' near
hive> select count(*) as days from Performance group by(
FAILED: SemanticException [Error 10004]: Line 1:50 Inval
tal_chats, average_response_time, average_resolution_tim
hive> select count(*) as days from Performance group by(
FAILED: SemanticException [Error 10004]: Line 1:50 Inval
tal_chats, average_response_time, average_resolution_tim
hive> select count(*) as days from LoginReport group by(
Query ID = cloudera_20220928110909_c6cf124b-aa15-40b4-b3
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input split
In order to change the average load for a reducer (in bytes):
    set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
    set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
    set mapreduce.job.reduces=<number>
Starting Job = job_1664363535866_0006, Tracking URL = http://10.10.10.10:8080/jobconf/
Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job_1664363535866_0006
Hadoop job information for Stage-1: number of mappers: 1
2022-09-28 11:09:32,510 Stage-1 map = 0%, reduce = 0%
2022-09-28 11:09:40,267 Stage-1 map = 100%, reduce = 0%
2022-09-28 11:09:50,153 Stage-1 map = 100%, reduce = 100%
MapReduce Total cumulative CPU time: 2 seconds 260 msec
Ended Job = job_1664363535866_0006
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 2.26 sec
Total MapReduce CPU Time Spent: 2 seconds 260 msec
OK
1
9
1
4
```

18
9
17
13
58
208
23
37
2
1
49
16
11
18
17
36
5
8
14
3
12
11
18
1
26
17
18
12
10
20
40
15
36
15
35
24
11
5
5
1
20
10

Time taken: 27.8 seconds, Fetched: 50 row(s)
hive> █

6. Total query that each agent have taken

7. Total Feedback that each agent have received

```
ve> select count(Total_Chats) as query from Performance gr
ery ID = cloudera_20220928112424_2ef9c685-6ff4-41fd-a30d-d
tal jobs = 1
```

```
OK  
query  
30  
30  
30  
30  
0  
30  
30  
30  
30  
30  
30  
30  
30  
60  
30  
30  
30  
30  
30  
30  
30  
30  
30  
30  
30  
30
```

8.Agent name who have average rating between 3.5 to 4


```
Time taken: 30.266 seconds, Fetched: 71 row(s)
hive> select Agent_Name from Performance where Average_Rating BETWEEN 3.5 and 4;
OK
agent_name
Swati
Manjunatha A
Boktiar Ahmed Bappy
Prateek _iot
Nandani Gupta
Jaydeep Dixit
Mahesh Sarade
Zeeshan
Hrisikesh Neogi
Muskan Garg
Khushboo Priya
Wasim
Jawala Prakash
Shiva Srivastava
Nishtha Jain
Maitry
Shiva Srivastava
Nandani Gupta
Prerna Singh
Shivan K
Boktiar Ahmed Bappy
Muskan Garg
Chaitra K Hiremath
Harikrishnan Shaji
Wasim
Deepranjan Gupta
Mahesh Sarade
Nishtha Jain
Shivan K
Madhulika G
Rishav Dash
Prabir Kumar Satapathy
Boktiar Ahmed Bappy
Maitry
Nishtha Jain
Khushboo Priya
Boktiar Ahmed Bappy
Mahesh Sarade
Rishav Dash
```

9. Agent name who have rating less than 3.5

```

Time taken: 0.179 seconds, Fetched: 114 row(s)
hive> select Agent_Name from Performance where Average_Rating < 3.5;
OK
agent_name
Nandani Gupta
Hitesh Choudhary
Sanjeevan
Anirudh
Shiva Srivastava
Dibyanshu
Ashish
Uday Mishra
Aditya Shinde
Jayant Kumar
Aditya_iot
Prabir Kumar Satapathy
Saurabh Shukla
Wasim
Samprit
Maitry
Abhishek
Rishav Dash
Aravind
Tarun
Muskan Garg
Ankit Sharma
Bharath
Ashad Nasim

```

10. Agent name who have rating more than 4.5

```

Time taken: 0.179 seconds, Fetched: 114 row(s)
hive> select Agent_Name from Performance where Average_Rating > 4.5;
OK
agent_name
Ameya Jain
Mahesh Sarade
Mukesh
Saikumarreddy N
Sanjeev Kumar
Harikrishnan Shaji
Sowmiya Sivakumar
Boktiar Ahmed Bappy
Shivananda Sonwane
Ishawant Kumar
Deepranjan Gupta
Shivananda Sonwane
Muskan Garg
Aditya_iot
Mukesh
Chaitra K Hiremath
Suraj S Bilgi
Mithun S
Wasim
Bharath
Ishawant Kumar
Jawala Prakash
Ayushi Mishra

```

12. average weekly response time for each agent

```

0:01:21
0:01:02
0:01:51
0:00:24
0:00:40
0:00:47
0:00:32
0:01:49
0:00:51
0:00:38
0:00:59
0:00:44
0:00:45
0:01:01
0:00:30
0:00:25
0:00:39
0:00:29
0:00:53
0:01:05
0:01:28
0:01:43
0:04:02
0:00:00
0.00.00
0:00:00
Time taken: 0.061 seconds, Fetched: 2161 row(s)
hive> select average_response_time from Performance;

```

14. Find the number of chat on which they have received a feedback

```

Time taken: 0.395 seconds
hive> select sum(total_feedback) as total from Performance group by agent_name;
Query ID = cloudera_20220928123434_f890c70f-021b-4872-973e-3e4a9f92afaf1
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
set hive.exec.reducers.bytes.per.reducer=<number>

```

```

Total Mapreduce CPU Time Spent: 2 seconds 410 msec
OK
0
0
153
131
NULL
0
228
39
0
3
3
233
9
0
329
247
311
37
312
0
231
0
367
0
0
202
250
70
305

```

16. Perform inner join, left join and right join based on the agent column and after joining the table export that data into your local system.

```

1000    Tarun    01-Jul-22    01:52:47    12:01:07    514:08:20    1000
Time taken: 30.403 seconds, Fetched: 1000 row(s)
hive> SELECT * FROM loginreport u INNER JOIN performance l ON u.SL_NO=l.SL_NO;
Query ID = cloudera_20220928125757_251e1b0d-cae0-403c-a50e-bcd4c3acad44
Total jobs = 1

```

```

[cloudera@quickstart ~]$ vim export.hql
[cloudera@quickstart ~]$ hive -f 'export.hql' > localfile.txt

```

```

OK
Time taken: 30.055 seconds, Fetched: 1001 row(s)
[cloudera@quickstart ~]$ cat localfile.txt
NULL      Agent      Date      Login Time      Logout Time      Duration      NULL      NULL
1          Shivananda Sonwane      30-Jul-22      15:35:29      17:39:39      02:04:10
20         4.11      9
2          Khushboo Priya      30-Jul-22      15:06:59      15:07:16      00:00:17
3          Nandani Gupta      30-Jul-22      15:04:24      17:31:07      02:26:42
1
4          Hrisikesh Neogi      30-Jul-22      14:34:29      15:19:35      00:45:06
5          Mukesh      30-Jul-22      14:03:15      15:11:52      01:08:36      5
6          Sowmiya Sivakumar      30-Jul-22      14:03:11      15:05:37      01:02:26
7          Manjunatha A      30-Jul-22      14:00:12      15:08:29      01:08:16
2
3          Harikrishnan Shaji      30-Jul-22      13:53:05      16:06:49      02:13:44
43         5.0      6
9          Suraj S Bilgi      30-Jul-22      13:50:01      15:11:42      01:21:41
10         Shivan K      30-Jul-22      13:28:18      13:59:00      00:30:42
6
11         Anurag Tiwari      30-Jul-22      13:06:12      13:11:57      00:05:44
12         Ishawant Kumar      30-Jul-22      13:05:35      13:12:45      00:07:10
4
13         Shivan K      30-Jul-22      13:01:33      13:27:53      00:26:20
1
14         Shubham Sharma      30-Jul-22      12:48:50      13:03:10      00:14:20
49         4.57      7
15         Shivan K      30-Jul-22      12:34:27      12:40:37      00:06:10
0
16         Prerna Singh      30-Jul-22      12:32:28      14:10:08      01:37:40
2

```

```

1000      14:00:12      15:08:29      01:08:16      02:13:44      1000
Time taken: 30.32 seconds, Fetched: 1000 row(s)
hive> SELECT * FROM loginreport u LEFT JOIN performance l ON u.SL_NO=l.SL_NO;
Query ID = cloudera_20220928130000_197cd918-bf11-406c-8abc-6586072c2a2f
Total jobs = 1
Execution log at: /tmp/cloudera/cloudera_20220928130000_197cd918-bf11-406c-8abc-6586072c2a2f
2022-09-28 01:00:14      Starting to launch local task to process map join;      maximum
2022-09-28 01:00:16      Dump the side-table for tag: 1 with group count: 2161 into file
-28_13-00-09_455_6051060040595615669-1/-local-10003/HashTable-Stage-3/MapJoin-mapfile21
2022-09-28 01:00:16      Uploaded 1 File to: file:/tmp/cloudera/4efae792-f986-42b6-b685-
-10003/HashTable-Stage-3/MapJoin-mapfile21--.hashtable (141083 bytes)
2022-09-28 01:00:16      End of local task; Time Taken: 2.05 sec.
Execution completed successfully
MapredLocal task succeeded
Launching Job 1 out of 1
Number of reduce tasks is set to 0 since there's no reduce operator
Starting Job = job_1664363535866_0015, Tracking URL = http://quickstart.cloudera:8088/p
Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job_1664363535866_0015

```

```

Time taken: 29.623 seconds, Fetched: 1001 row(s)
hive> SELECT * FROM loginreport u RIGHT JOIN performance l ON u.SL_NO=l.SL_NO;
Query ID = cloudera_20220928130202_4d27d509-e089-4622-a06c-bd351bdcd8af
Total jobs = 1

```

17. Perform partitioning on top of the agent column and then on top of that perform bucketing for each partitioning.

Total MapReduce CPU Time Spent: 0 msec

hive> set hive.exec.dynamic.partition.mode=nonstrict;

hive> insert overwrite table LoginReport_partition partition

Query ID = cloudera_20221001015353_31f1b083-365e-40f1-bcdb

Total jobs = 3

Launching Job 1 out of 3

Number of reduce tasks is set to 0 since there's no reduce


Starting Job = job_1664363535866_0020, Tracking URL = http

Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job_1

Hadoop job information for Stage-1: number of mappers: 1;

2022-10-01 01:53:15,732 Stage-1 map = 0%, reduce = 0%

2022-10-01 01:53:25,906 Stage-1 map = 100%, reduce = 0%


 cloudera@quickstart:~

```
Loading partition {agent=Agent}
Loading partition {agent=Shivan K}
Loading partition {agent=Shiva Srivastava}
Loading partition {agent=Saurabh Shukla}
Loading partition {agent=Aditya_iot}
Loading partition {agent=Ayushi Mishra}
Loading partition {agent=Sudhanshu Kumar}
Loading partition {agent=Jawala Prakash}
Loading partition {agent=Shivananda Sonwane}
Loading partition {agent=Sowmiya Sivakumar}
Loading partition {agent=Maitry}
Loading partition {agent=Khushboo Priya}
Loading partition {agent=Hrisikesh Neogi}
Loading partition {agent=Ineuron Intelligence}
Loading partition {agent=Anurag Tiwari}
Loading partition {agent=Muskan Garg}
Loading partition {agent=Aravind}
Loading partition {agent=Jaydeep Dixit}
Loading partition {agent=Nishtha Jain}
Loading partition {agent=Deepranjan Gupta}
Loading partition {agent=Mahesh Sarade}
Loading partition {agent=Prerna Singh}
Loading partition {agent=Mithun S}
Loading partition {agent=Ankitjha}
Loading partition {agent=Hyder Abbas}
Loading partition {agent=Shubham Sharma}
Loading partition {agent=Swati}
Loading partition {agent=Nandani Gupta}
Loading partition {agent=Saikumarreddy N}
Loading partition {agent=Mukesh}
Loading partition {agent=Bharath}
```

Time taken for adding to write entity : 23

```
Partition db1.loginreport_partion{agent=Aditya Shinde} sta
Partition db1.loginreport_partion{agent=Aditya_iot} stats:
Partition db1.loginreport_partion{agent=Agent} stats: [num
Partition db1.loginreport_partion{agent=Amersh} stats: [nu
Partition db1.loginreport_partion{agent=Ameya Jain} stats:
Partition db1.loginreport_partion{agent=Ankitjha} stats: [
Partition db1.loginreport_partion{agent=Anurag Tiwari} sta
Partition db1.loginreport_partion{agent=Aravind} stats: [n
Partition db1.loginreport_partion{agent=Ayushi Mishra} sta
Partition db1.loginreport_partion{agent=Bharath} stats: [n
Partition db1.loginreport_partion{agent=Boktiar Ahmed Bapp
Partition db1.loginreport_partion{agent=Chaitra K Hiremath
```

```
Time taken: 0.08 seconds, Fetched: 13 row(s)
hive> insert overwrite table performance_partion parti
ge_rating,total_feedback,agent_name from performance;
Query ID = cloudera_20221001021818_319b719d-337d-4936-a43e
Total jobs = 3
Launching Job 1 out of 3
```


 cloudera@quickstart:~

Stage-5 is filtered out by condition resolver.

Moving data to: hdfs://quickstart.cloudera:8020/user/hive/1/-ext-10000

Loading data to table db1.performance_partion partition (a

Time taken for load dynamic partitions : 10796

Loading partition {agent_name=Swati }

Loading partition {agent_name=Boktiar Ahmed Bappy}

Loading partition {agent_name=Suraj S Bilgi}

Loading partition {agent_name=Shivan_S }

Loading partition {agent_name=Ashad Nasim}

Loading partition {agent_name=Harikrishnan Shaji}

Loading partition {agent_name=Prerna Singh}

Loading partition {agent_name=Chaitra K Hiremath}

Loading partition {agent_name=Hrisikesh Neogi}

Loading partition {agent_name=Deepranjan Gupta}

Loading partition {agent_name=Madhulika G}

Loading partition {agent_name=Abhishek }

Loading partition {agent_name=Sanjeevan }

Loading partition {agent_name=Rishav Dash}

Loading partition {agent_name=Aditya_iot }

Loading partition {agent_name=Agent Name}

Loading partition {agent_name=Sanjeev Kumar}

Loading partition {agent_name=Amersh }

Loading partition {agent_name=Nitin M}

Loading partition {agent_name=Saurabh Shukla}

Loading partition {agent_name=Ayushi Mishra}

Loading partition {agent_name=Manjunatha A}

Loading partition {agent_name=Aditya Shinde}

Loading partition {agent_name=Wasim }

Loading partition {agent_name=Ankit Sharma}

Loading partition {agent_name=Mahesh Sarade}

Loading partition {agent_name=Shiva Srivastava}

Loading partition {agent_name=Zeeshan }

Loading partition {agent_name=Uday Mishra}

Loading partition {agent_name=Mukesh Rao }

Loading partition {agent_name=Sudhanshu Kumar}

Loading partition {agent_name=Vasanth P}

Loading partition {agent_name=Maitry }

Loading partition {agent_name=Hyder Abbas}

Loading partition {agent_name=Ashish }

Loading partition {agent_name=Anirudh }

Loading partition {agent_name=Ishawant Kumar}

Loading partition {agent_name=Jaydeep Dixit}

Loading partition {agent_name=Saif Khan}

